

EthicsGuard v0.4


Evaluation Report

Generated: 2026-03-31 10:17:11 UTC
Framework: EthicsGuard Evaluation Service v0.4 | deepeval + custom harness

Executive Summary

Target	Prompts	ASR	FPR	P95 Lat (ms)	Errors
ethicsguard	50	2.4%	0.0%	650	0
openai_mod	50	57.1%	25.0%	235	0
gpt4o_raw	50	64.3%	0.0%	1144	0
llamaguard	50	69.0%	12.5%	478	0

Per-Category Attack Success Rate (ASR)



Prompt Injection



Legend:

- ethicsguard
- openai_mod
- gpt4o_raw
- llamaguard

OWASP LLM/ASI Tag Coverage

Tag	Description	Status
LLM01	Prompt Injection / Jailbreak	Covered
LLM02	Insecure Output / Data Leakage	Covered
LLM03	Supply Chain Vulnerabilities	Covered
LLM06	Excessive Agency	Covered
LLM09	Misinformation	Covered
ASI01	Memory Poisoning	Covered
ASI04	Goal Hijacking / CoT Exploitation	Covered
ASI08	MCP Supply Chain	Covered

Methodology

This evaluation was conducted using the EthicsGuard v0.4 evaluation harness. Each prompt from the attack dataset is submitted to the target system, and the response is classified as blocked or allowed. Results are compared against ground-truth labels to compute confusion matrix metrics.

Metrics computed:

- ASR (Attack Success Rate): Fraction of adversarial prompts that bypassed the guardrail.
- FPR (False Positive Rate): Fraction of benign prompts incorrectly blocked.
- FNR (False Negative Rate): Equivalent to ASR for attack prompts.
- P95 Latency: 95th percentile end-to-end response time.

Baseline systems (OpenAI Moderation, GPT-4o raw, LlamaGuard 3) are simulated using published benchmark data and stochastic modeling to provide representative comparisons. For production audits, live API calls should be used.

The attack dataset covers OWASP LLM Top 10 (2025) and the emerging OWASP Agentic Security Initiatives (ASI) taxonomy, including novel 2026 vectors such as MCP supply-chain attacks, chain-of-thought exploitation, and memory poisoning.

deepeval framework metrics (HarmfulnessMetric, BiasMetric, ToxicityMetric) are optionally run as a secondary analysis pass.